


# Exploring Stochastic Gradient Descent and its Modifications

Kevin Kleiner

Algorithm Interest Group on March 22, 2021



# Model Fitting and Optimization


- From linear models to neural networks, here are some problems which call for model fitting

Collected training data

1. Crystal structures and respective optical properties
2. Time series of relative positions of celestial objects
3. Time series of many companies' stock prices

Optimized model


1. Structure-property relationships to inform design
2. Planetary trajectories to inform mapping
3. Stock predictions to inform investing



Which model fitting problems have you encountered in your work?

# Model Fitting and Optimization

- All these models involve “weights” that should correctly map the inputs to an output
- We need a rigorous, systematic method to determine the “best possible” model weights



Animation from: <https://hackernoon.com/visualizing-linear-regression-with-pytorch-9261f49edb09>

# Optimize Model Weights with Gradient Descent

- Introducing the most popular approach,

$$w_{t+1} = w_t - \eta \nabla_{w_t} L(\{x_i, y_i\}; w_t)$$

Weights at iteration  $t + 1$

Weights at iteration  $t$

Step size:  $0 < \eta < 1$


Loss function

Set of all observations/training data

$$L_{LS}(\{x_i, y_i\}; w_t) \equiv \sum_i (\hat{y}_i(x_i; w_t) - y_i)^2$$

Least squares loss function


Model-predicted  $y$  using  $x$  and  $w_t$



# Each Gradient Evaluation can be Expensive

- Say we want to fit a linear model  $\hat{y} = w_0 + w_1x_1 + w_2x_2 + \cdots + w_Nx_N$  to  $N \sim 10^6$  data points

$$\nabla_{w_t} L = \nabla_{w_t}(\hat{y}_1(x_1; w_t) - y_1)^2 + \nabla_{w_t}(\hat{y}_2(x_2; w_t) - y_2)^2 + \cdots + \nabla_{w_t}(\hat{y}_N(x_N; w_t) - y_N)^2$$



Derivatives  
with respect to  
all  $N$  weights

Gradient operates  
on all  $N$  residual  
functions

- How can the weights be updated more frequently, while still converging to the correct loss minimum?
- Consider a slight modification to improve scaling and see how the optimization behavior changes

# Stochastic Gradient Descent Enables Faster Evaluations

- Instead of updating weights with  $\nabla_{w_t} L(\{x_i, y_i\}; w_t)$ , update them with  $\nabla_{w_t} L_i(x_i, y_i; w_t)$
- This simplified loss function only depends on one randomly chosen observation  $(x_i, y_i)$

$$w_{t+1} = w_t - \eta \nabla_{w_t} L_i(x_i, y_i; w_t)$$

$$L_{LS,i}(x_i, y_i; w_t) \equiv (\hat{y}_i(x_i; w_t) - y_i)^2$$

Each red update is much, much faster than each black update, and both paths should converge to roughly the same point




Image from: David Macedo PhD Thesis 2017

# SGD Updates are too Sensitive to Individual Data

- So weight optimization trajectory becomes very noisy, especially after the beginning
- We now want to reduce the SGD noise without returning to GD scaling

- Replace  $\nabla_{w_t} L_i(x_i, y_i; w_t)$  with  $\nabla_{w_t} L(\{x_i, y_i\}_{MB}; w_t)$
- $\{x_i, y_i\}_{MB}$  is a random mini-batch of data points

## Mini-Batch Stochastic Gradient Descent




Image from: <https://towardsdatascience.com/gradient-descent-algorithm-and-its-variants-10f652806a3>

# **Break: Demonstrate Previous Methods in Live Demo**

Using height-weight data from Kaggle:

<https://www.kaggle.com/mustafaali96/weight-height>

# Modifications Beyond Mini-Batch SGD

- Further modifications become problem-dependent, which requires knowing the data well
- Suppose the weight optimization still experiences difficulties:
  1. Convergence to an undesired local minimum - depends on loss function shape
  2. Noisy zig-zag updates - depends on redundancies or bias in the data
    - A common solution is introducing “momentum”



(a) SGD without momentum



(b) SGD with momentum

# Mini-Batch SGD with Momentum

- Update not only the weights, but also a “weight velocity”
- Analogous to a velocity-dependent damping force in Newton’s laws

$$w_{t+1} = w_t - v_t$$

$$v_t = \alpha v_{t-1} + \eta \nabla_{w_t} L(\{x_i, y_i\}_{MB}; w_t)$$

1. Speeds up only the updates heading towards the minimum
2. Pushes updates out of possible traps



(a) SGD without momentum



(b) SGD with momentum

# Visual Schematic of SGD with Momentum

$$v_t = \alpha v_{t-1} + \eta \nabla_{w_t} L(\{x_i, y_i\}_{MB}; w_t)$$

$$w_{t+1} = w_t - v_t$$




Image from: <https://towardsdatascience.com/a-visual-explanation-of-gradient-descent-methods-momentum-adagrad-rmsprop-adam-f898b102325c>

# Visual Schematic of SGD with Momentum Cont.





Image from: <https://towardsdatascience.com/a-visual-explanation-of-gradient-descent-methods-momentum-adagrad-rmsprop-adam-f898b102325c>

# Sometimes SGD with Momentum Still Isn't Enough

- If the loss function has a deep valley, our previous methods don't work


Example: logistic regression on the noisy moons dataset in scikit-learn



Let's see why these other methods help

Animation from: <http://www.denizyuret.com/2015/03/alec-radfords-animations-for.html>

# Including Adaptive Step Sizes




$$L \rightarrow L(\{x_i, y_i\}_{MB}; w_t - \alpha v_{t-1})$$

$$\eta_0 \rightarrow \frac{\eta_0}{\sqrt{\sum_{\tilde{t}=1}^t (\nabla_{w_{\tilde{t}}} L)^2 + \epsilon}}$$

Replaces Adagrad sum  
with weighted average

Animation from: <http://www.denizyuret.com/2015/03/alec-radfords-animations-for.html>

# Connecting Back to Model Accuracy




Animation from: <http://www.denizyuret.com/2015/03/alec-radfords-animations-for.html>

# Conclusions: Zoo of Optimization Methods

| Method                    | Advantage                           | Disadvantage                      |
|---------------------------|-------------------------------------|-----------------------------------|
| <b>GD</b>                 | Least noisy updates                 | Expensive gradient evaluations    |
| <b>SGD</b>                | Cheap gradient evaluations          | Most noisy updates                |
| <b>Mini-Batch SGD</b>     | Cheap gradient evaluations          | Moderately noisy updates          |
| <b>SGD + Momentum</b>     | More traversal in desired direction | Unreliable in non-convex terrains |
| <b>Adagrad and Beyond</b> | Self-corrected step sizes           | Step sizes vanish after a while   |

- When confronted with a model fitting problem with lots of data, choose an optimization method by assessing:
  1. Whether it can reliably converge to a loss minimum
  2. How long it takes to converge
  3. What path it takes through the loss function terrain

# Short Overview Video on Optimizers



“Optimizers - EXPLAINED!” by CodeEmporium